

```

# Allus definitions.
# You may want to put all your additions into a separate file like
# ~/.bash_aliases, instead of adding them here directly.
# See /usr/share/doc/bash-doc/examples in the bash-doc package.

if [ -f ~/.bash_aliases ] then
. ~/.bash_aliases
fi

# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
if [ -f /usr/share/bash-completion/bash_completion ] then
. /usr/share/bash-completion/bash_completion
elif [ -f /etc/bash_completion ] then
. /etc/bash_completion
fi
fi

export JAVA_HOME=/usr/lib/jvm/jdk1.8.0_202
export PATH=/usr/local/java/jdk/bin
export HADOOP_HOME=/usr/hadoop
export PATH=/usr/hadoop/bin
export PATH=/usr/sbin
export HADOOP_MAPRED_HOME=/usr/hadoop-mapred
export YARN_HOME=/usr/hadoop-yarn
export HADOOP_CONF_DIR=/usr/hadoop/etc/hadoop
export HADOOP_STREAMING=/usr/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.4.0.jar
export HADOOP_LOG_DIR=/usr/hadoop/logs
export PS4=$PWD_TYPE-$HOSTNAME
export HADOOP_COMMON_LIB_NATIVE_DIR=/usr/hadoop/lib/native
export HADOOP_OPTS="-Djava.library.path=/usr/hadoop/lib/native"

```

[illegible]

\$nano \$HADOOP_HOME/etc/hadoop/core-site.xml

```

File Actions Edit View Help
GNU nano 7.2 /home/hac
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the license is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the license for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>fs.defaultFS</name>
<value>hdfs://localhost:9000</value> </property>
</property>
<name>hadoop.proxyuser.dataflair.groups</name> <value>*</value>
</property>
<name>hadoop.proxyuser.dataflair.hosts</name> <value>*</value>
</property>
<name>hadoop.proxyuser.server.hosts</name> <value>*</value>
</property>
<name>hadoop.proxyuser.server.groups</name> <value>*</value>
</property>
</configuration>

```

\$nano \$HADOOP_HOME/etc/hadoop/hdfs-site.xml

```

File Actions Edit View Help
GNU nano 7.2 /home/hadoop/hadoop/etc/hadoop
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the license is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the license for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<name>dfs.name.dir</name>
<value>file:///home/hadoop/hadoopdata/hdfs/namenode</value>
</property>
<name>dfs.datanode.data.dir</name>
<value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
</property>
</configuration>

```

\$nano \$HADOOP_HOME/etc/hadoop/mapred-site.xml

```

hadoop@kali: ~
File Actions Edit View Help
GNU nano 2.2 /home/hadoop/hadoop/etc/hadoop/mapred-site.xml

<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>mapreduce.framework.name</name> <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.application.classpath</name>
    <value>${HADOOP_MAPRED_HOME}/share/hadoop/mapreduce/*:${HADOOP_MAPRED_HOME}/share/hadoop/mapreduce/lib/*</value>
  </property>
</configuration>

```

\$nano \$HADOOP_HOME/etc/hadoop/yarn-site.xml

```

hadoop@kali: ~
File Actions Edit View Help
GNU nano 2.2 /home/hadoop/hadoop/etc/hadoop/yarn-site.xml
<?xml version="1.0"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->
<configuration>
<!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PREPENDED_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>
  </property>
</configuration>

```

\$ start-all.sh

```

hadoop@kali: ~
File Actions Edit View Help
(hadoop@kali)~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 18 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [kali]
Picked up _JAVA_OPTIONS: -Dont.useSystemAAFontSettings-on -Dswing.text=true
2024-09-11 04:58:16,429 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting resourcemanager
Starting node managers

```

\$ jps

```
(kannapp@kali) ~$ jps
Picked up _JAVA_OPTIONS: -Dawt.useSystemAAFontSettings=on -Dswing.aatext=true
14436 NodeManager
16772 Jps
13838 SecondaryNameNode
16331 ResourceManager
13597 DataNode
13471 NameNode
(kannapp@kali) ~$
```

localhost:9870

The screenshot shows a web browser window displaying the Apache Hadoop 3.3.6 WebHDFS interface. The page title is "SecondaryNameNode info". The interface includes a navigation bar with "Hadoop" and "Overview" tabs. The "Overview" tab is active, showing a table of Hadoop metrics.

Metric	Value
Version	3.3.6
Compiled	2023-06-18T06:22Z by ubuntu from (HEAD detached at release-3.3.6-RC1)
NameNode Address	localhost:9000
Started	Wed Aug 14 21:51:32 -0400 2024
Last Checkpoint	Never
Checkpoint Period	3600 seconds
Checkpoint Transmitters	1000000

Below the table, there are sections for "Checkpoint Image URI" and "Checkpoint Editing URI", both showing a default value: `file:///tmp/hadoop-kannapp/dfs/secondary/`.

The bottom of the screenshot shows the Windows taskbar with the time 10:08 AM on 15-08-2024.

localhost:8088

The screenshot shows the Apache Hadoop 3.3.6 web interface. The top navigation bar includes links for Home, Clusters, Nodes, Scheduler, and Applications. The main content area is titled 'All Applications' and displays several metrics and a table.

Cluster Metrics

App Submitted	App Pending	App Running	App Completed	Taskset Running	Local Resources	Total Resources
0	0	0	0	0	<memory:0 B, vCores:0>	<memory:0 B, vCores:0>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Default
0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	Memory MB (MB=MB, vCores)	<memory:1GB, vCores:1>	<memory:8192, vCores:4>

Applications Table

ID	State	Name	Application Type	Application Type	Queue	Application Priority	StartTime	LastestTime	FinishTime	State	FinalState	Running Containers	Allocated CPU V-Cores	Allocated Memory MB	Allocated GPUs
No data available in table															

Showing 0 to 0 of 0 entries