

EXP NO: 5

INSTALLATION OF HIVE

\$wget <https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz>



\$tar xvfz apache-hive-3.1.2-bin.tar.gz




\$mv apache-hive-3.1.2-bin hive

\$nano ~/.bashrc

\$source ~/.bashrc

\$cd hive

\$cp conf/hive-default.xml.template conf/hive-site.xml



\$nano conf/hive-site.xml


```
>show databases;
```

[illegible]

```
hive> CREATE DATABASE financials;
OK
Time taken: 0.591 seconds
```

```
hive> use financials;
OK
Time taken: 0.229 seconds
```

```
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 2.536 seconds
```

```

hive> INSERT INTO Finance_Table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Copy 10 = hadoop-20240910011318-799e97fa-d8cc-4d28-a341-1708acc000a9
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes-per-reducer=number
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=number
In order to set a constant number of reducers:
  set mapreduce.job.reduces=number
Launching Job = job_1725209231018_0001, Tracking URL = http://kali:8080/progress/application_1725209231018_0001/
kill Command = /home/hadoop/hadoop/bin/mapred job -kill job_1725209231018_0001
hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
  2024-09-10 01:1A:33:477 Stage-1 map = 0%, reduce = 0%
  2024-09-10 01:15:37:392 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 28.45 sec
  2024-09-10 01:15:37:443 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 33.61 sec
MapReduce Total cumulative CPU time: 33 seconds 818 msec
Finished Job = job_1725209231018_0001
Stage-4 is selected by condition resolver.
Stage-5 is filtered out by condition resolver.
Stage-6 is filtered out by condition resolver.
Writing data to directory hdfs://localhost:9000/user/kive/warehouse/financials.db/finance_table/hive-staging_hive_2024-09-10_01-13-18_205_2162523115504052-1/-ext-190008
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-31 Map: 1 Reducer: 1 Cumulative CPU: 33.61 sec HDFS Read: 1566B HDFS Write: 291 SUCCESS
Total MapReduce CPU Time Spent: 33 seconds 818 msec
OK
Time taken: 151.732 seconds

```

> CREATE VIEW myview AS SELECT name, id FROM finance_table;

```
hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;  
OK  
Time taken: 0.82 seconds
```

> SELECT * FROM myview;

```
hive> SELECT * FROM myview;  
OK  
Alice 1  
Bob 2  
Charlie 3  
Time taken: 0.716 seconds, Fetched: 3 row(s)  
hive> 
```

> DESCRIBE finance_table;

```
hive> DESCRIBE finance_table;  
OK  
id int  
name string  
Time taken: 0.348 seconds, Fetched: 2 row(s)
```

> ALTER TABLE finance_table ADD COLUMNS (age INT);

>quit;

```
hive> ALTER TABLE finance_table ADD COLUMNS (age INT);  
OK  
Time taken: 0.457 seconds  
hive> quit;
```